

Machine Learning (HWS22)

Assignment 4: Latent Variable Models

The archive provided to you contains this assignment description, a dataset in a binary format, as well as Python code fragments for you to complete. Comments and documentation in the code provide further information.

It suffices to fill out the “holes” that are marked in the code fragments provided to you, but feel free to modify the code to your liking. You need to stick with Python though.

Provide a single ZIP archive with name `ml22-<assignment number>-<your ILIAS login>.zip`. The archive needs to contain:

- A **single PDF report** that contains answers to the tasks specified in the assignment. Do not simply convert your Jupyter notebook to a PDF! Write a separate document, stay focused and brief. **Use at most 10 single-column pages.**
- All the **code** that you created and used in its original format.
- A PDF document that renders your Jupyter notebook with all figures. (If you don’t use Jupyter, then you obviously do not need to provide this.)

You need to adhere to the above guidelines in your submission, otherwise we may grade your solution as a **FAIL**.

Generally, your report should

- include a high-level description of your approach and helpful figures,
- be self-explanatory (i.e., refer to code *only* for implementation-only tasks),
- follow standard scientific practice,
- include appropriate references if you used additional sources or material,
- not include any hand-written notes,
- label all figures/tables and refer to figures/tables via their labels,
- use one section per task and one subsection per subtask, each numbered with the (sub)task numbers from the assignment sheet.

Your report will be downgraded if you do not follow these points (e.g., you can’t get **EXCELLENT**).

Hand-in your solution via ILIAS until the date specified there. This is a hard deadline.

1 Probabilistic PCA

- a) Study the PPCA generator provided to you (`ppca_gen`). Generate and plot the `toy_ppca` dataset. What is shown in the plot? Vary the amount of noise (σ^2), replot, and describe how the data distribution changes (visually) depending on σ^2 .
- b) Implement MLE for PPCA by completing `ppca_mle`. Only use standard matrix/vector operations and the `svd` function. If we fit the PPCA model with $L = 2$ on `toy_ppca`, we obtain $\hat{\sigma}_{\text{MLE}}^2 = 0$. Why is this?
- c) Implement the computation of the conditional negative log-likelihood of a given dataset for a given PPCA model. To do so, complete `ppca_nll`.
- d) Load the `secret_ppca` dataset, which was produced by sampling a PPCA model. Discover its secret: How many latent variables have been used (L)? Do this by
 - (i) Studying the scree plot.
 - (ii) Using validation data.

Do the results agree?

2 Gaussian Mixture Models

- a) Study the GMM generator provided to you (`gmm_gen`). Generate and plot the `toy_ppca` dataset. Describe what is shown in the plot.
- b) Compute a K -Means clustering of `toy_ppca` using $K = 5$ (code provided). Plot the resulting clustering and discuss. Did you expect this result?
- c) Implement the E step and the M step of fitting a GMM with MLE by completing `gmm_e` and `gmm_m`. Only use standard matrix/vector operations and (optionally) the functions mentioned in the hints below.

Hint (E step). Start by computing and verifying \mathbf{F} ($[\mathbf{F}]_{ik} = f_k(\mathbf{x}_i)$ in lecture notation). You can compute the density of each data point (row) in a dataset (\mathbf{X}) under multivariate normal distribution $\mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma})$ using:

```
dist = scipy.stats.multivariate_normal(mean=mu, cov=Sigma)
densities = dist.pdf(X)
```

Hint (M step). Start by computing and verifying $\boldsymbol{\pi}$, then the $\boldsymbol{\mu}_k$'s, then the $\boldsymbol{\Sigma}_k$'s. Perhaps helpful: NumPy's covariance function `np.cov` supports weighted data points (set `aweights` argument accordingly; additionally set `dof=0`).

- d) Fit the GMM model with $K = 5$ using `gmm_fit` (provided). Assign each data point to its most likely component to obtain a clustering and plot the dataset as in subtask b). Compare the resulting clustering with the result of K -Means clustering of subtask b).
- e) Repeat subtask d) with $K = 4$ and with $K = 6$. In both cases, repeat multiple times. Discuss your findings.
- f) Optional: Load the `secret_gmm` dataset, which was produced by sampling a GMM model. Discover it's secret: How many components (L) have been used? Briefly justify your answer.